

Airbnb Toronto Data Science Project

Team 4

Deliverable 2 Report : Data Exploration and Data cleaning

Sheridan College

Date: October 28th 2021

Table of Contents

Purpose of Deliverable 2

Overview of the data cleaning process and relevance of EDA

Part I

Steps in the removal of features

Imputation of missing values in Categorical features

Outlier detection and transformation of data

Part II

Feature Engineering

Imputation of missing values in Numerical features

Dummy coding categorical variables

Part III

EDA of original (unclean) dataset

Visualization from Tableau of Clean (cleansed) dataset

Conclusion

Appendix

TEAM 4-DELIVERABLE 2 : Data cleaning and Exploratory Data Analysis Report

Purpose of Deliverable 2

The goal of deliverable 2 of our project was to prepare a clean Toronto Airbnb dataset and provide data visualizations that would enable stakeholders to understand why essential steps characteristic of data cleaning; such as the removal of features and imputation took place. As noted in our earlier report, we believed that features in the dataset that are predictive of price of stay per night and review ratings score would enable the Airbnb company president and relevant stakeholders such as Airbnb hosts to develop business strategies such as marketing to emphasize on these features when providing Airbnb service to existing and potential customers. We expect a dramatic increase in Airbnb revenue based on the features to be presented in the results of this study. We utilized the programming language, Python for the steps required in data cleansing and Tableau (version 2021.3) for the Exploratory data analysis (EDA).

The Toronto Airbnb dataset presented a challenge in the data cleaning process because of four main reasons. First, we believe that there were too many variables/features (74, including listing ID) in the dataset. Upon closer examination, we identified features such as 'listing_url' that were not relevant to our business objective. We have also noted some variables that was redundant when comparing to other variables included in the dataset. We also removed variables if they contained more than 40% missing values. The complexity of the dataset is provided in Figure A (see Appendix) – where total number of 19,343 host listings included in our dataset is visually shown, imposed under a Toronto map.

Overview of the data cleaning process and relevance of EDA

Exploratory analysis (EDA) of the original dataset (uncleaned) was first conducted through the use of Tableau. Graphs and plots to identify relevant features and the relationships between the features were generated (refer to Part III of this report). Secondly, as is common with real world data, several features contained missing values and as result, imputation methods were implemented in Python to replace the missing values. Thirdly, outliers were identified in the dataset using both Tableau and Python. We detected

both negative and positive skewness in several features of the dataset. The interquartile range (IQR) method was selected to remove the outliers.

Finally, using feature engineering, we created three features that we think is relevant to our business objective. One feature created in the dataset is ‘former city’ based on Toronto districts. There were more than 40 neighborhoods and we decided to consolidate each neighborhood to a city. In Python, we web scraped the Wikipedia page <https://en.wikipedia.org/wiki/Toronto>, under the section entitled “Neighborhoods” to collect all the districts and associated neighbourhoods in Toronto. Using this method, we were able to group the neighbourhoods by district. Python screenshot below displays a sample list of neighbourhoods grouped by Toronto district (i.e., ‘former_city’).

```
[5]: df2[['price', 'review_scores_rating', 'neighbourhood_cleansed', 'former_city']].sample(10)
```

| | price | review_scores_rating | neighbourhood_cleansed | former_city |
|-------|----------|----------------------|---------------------------------|---------------------|
| 13980 | \$250.00 | NaN | Cabbagetown-South St.James Town | Old City of Toronto |
| 13795 | \$75.00 | 60.0 | Leaside-Bennington | East York |
| 16291 | \$80.00 | NaN | Lawrence Park South | Old City of Toronto |
| 17805 | \$45.00 | NaN | Dorset Park | Scarborough |

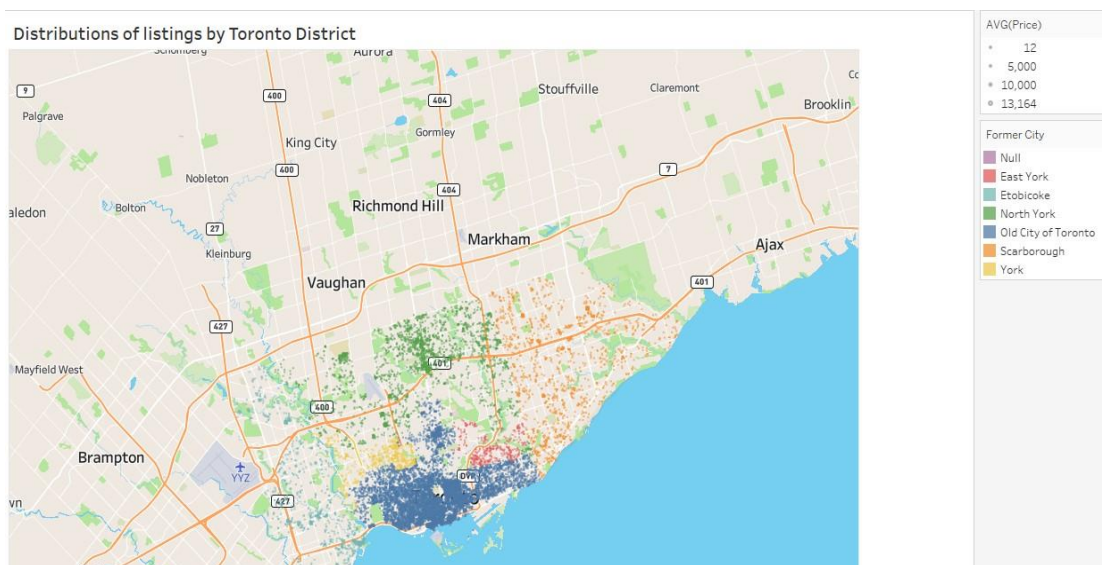
Python was used to generate the cleaned dataset that would be used for (1) EDA in Tableau for visualization and (2) for implementation of the machine learning model. We used Tableau for both uncleaned and cleaned datasets to see the effect of cleaning the data. Below, we provide in detail the steps taken to clean our Airbnb dataset using Python (Part I & Part II) and the necessary EDA in Tableau for the uncleaned dataset (Part III).

Part I

Steps in the removal of features

To begin, our EDA using Tableau to characterize price variation by neighborhood or review ratings score by neighborhood was difficult to interpret as there are more than 70 neighbourhoods in Toronto. Using Python, we web scraped the Wikipedia page of the city of Toronto to collect all the districts and relevant neighborhoods. As a result, we were able to group the neighbourhoods by district. A merge function was implemented in Python to combine the web scraped table with our uncleaned dataset. Using this, a new feature called 'Former City' was added into the Toronto Airbnb dataset. The new feature 'Former City', consists of six Toronto districts namely Etobicoke, North York, East York, Old City of Toronto, York and Scarborough. Each of these Toronto districts have more than five neighbourhoods, with Old City of Toronto containing the most neighborhoods. For example, East York covers Broadview North, Danforth East York, Leaside-Bennington, O'Connor-Parkview, Old East-York, Thorncliffe Park and Woodbine-Lumsden. In *Figure B* below, you can see how listings can be viewed better by district rather than by neighborhood as shown with the uncleaned dataset in Tableau in *Figure A* (see appendix).

Figure B. Listings by District in Toronto



In *Table 1* in the Appendix, we have highlighted features in the original or uncleaned dataset in terms of the basis of removal (i.e., irrelevant, too many missing values and redundancy). There were 74 features (Id being a primary key) in the uncleaned dataset. Features were removed because of redundancy and irrelevance to our business objective and if data has more than 40% missing values. We have also highlighted features that had NaN values.

We used Python to reduce the 74 features to a set of features that were relevant to our machine learning model. We first did a correlation plot with the target variable being price (see *Figure 2* in Appendix). In general, there were few features correlated to price, highlighting the importance of data cleaning. Our second step in data cleaning involved removing features not relevant to our business objective.

Moreover, several features had missing values, as can be seen from *Table 2* in Appendix. For example, 'host_neighbourhood' had 15,552 non-null values, and as a result, 'host_neighbourhood' had 3,791 missing values (i.e., $19,343 - 15,552 = 3,791$). We removed *id*, 'listing_url', 'scrape_id', 'last_scraped', 'name', 'description', 'neighbourhood_overview', 'picture_url', 'host_id', 'host_url', 'host_name', 'host_thumbnail_url'. Textual information such as description and url weblinks variables were also removed. Features such as 'minimum_nights_avg_ntm' and 'maximum_nights_avg_ntm' were removed because of redundancy in the data. That is, both features had duplicate values.

Thirdly, we removed features for which there were no values. Fourthly, we removed features that had more than 40% missing values. These features were 'host_about', 'neighbourhood_group_cleansed', 'bathrooms', 'calendar_updated', 'license'. The 40% missing value threshold is what most businesses use for removal of missing values in datasets. Refer to the figure below for the output from Python, in terms of percentage of missing values being higher than 40% for five features.

```
Columns with more than 40% missing values:
['host_about' 'neighbourhood_group_cleansed' 'bathrooms'
 'calendar_updated' 'license']
```

Imputation of missing values in Categorical features

As a result of the above steps of data cleaning, we had 37 features in our dataset. There were nine categorical and 28 numeric features. Refer to *Table 3* in Appendix for profile of features with data type. To replace NaN values in the categorical features with data type of object, we imported the SimpleImputer from Sklearn Python library to impute missing values with most frequent. The SimpleImputer is a scikit-learn class. We implemented the ‘frequent’ strategy from SimpleImputer to impute the missing values. Host response time, host response rate and host acceptance rate were categorical features for which NaN values were imputed.

```
df.head(3)
```

| | room_type | host_since | host_response_time | host_is_superhost | host_acceptance_rate | host_response_rate |
|---|-----------------|------------|--------------------|-------------------|----------------------|--------------------|
| 0 | Entire home/apt | 2008-08-08 | NaN | f | NaN | NaN |
| 1 | Entire home/apt | 2011-06-07 | NaN | f | NaN | NaN |
| 2 | Entire home/apt | 2012-06-01 | NaN | t | 100% | NaN |

```
df.head(3)
```

| | room_type | host_since | host_response_time | host_is_superhost | host_acceptance_rate | host_response_rate |
|---|-----------------|------------|--------------------|-------------------|----------------------|--------------------|
| 0 | Entire home/apt | 2008-08-08 | within an hour | f | 100% | 100% |
| 1 | Entire home/apt | 2011-06-07 | within an hour | f | 100% | 100% |
| 2 | Entire home/apt | 2012-06-01 | within an hour | t | 100% | 100% |

Outlier detection and transformation of data

As can be seen from the box plot and histogram distributions generated in Python and for the selected features in Tableau, displayed in *Figure 4* in appendix; most of the data contained within features was positively or negatively skewed. Several outliers were present within each feature. The outliers were noted from use of describe function and confirmed from the box plots and visualization of the distributions. For example, in the unclean or original dataset; for price, the maximum value is \$13,164 but the average price is \$141.28 and 75% are below \$150. Similarly, for ‘minimum_nights’, the maximum value is 1125, but the average is 10 and 75% below 5. For both bedrooms and beds, mean was 1.39 and 1.63 respectively;

maximum being 16 and 17. Outliers were detected in the features: *'price'*, *'accommodates'*, *'host_total_listings_count'*, *'reviews_per_month'*, *'bedrooms'*, *'beds'*, *'minimum_nights'*, *'maximum_nights'*, *'number_of_reviews'*, *'number_of_review_ltm'*, *'review_scores_accuracy'*, *'review_scores_cleanliness'*, *'review_scores_checkin'*, *'review_scores_communication'*, *'review_scores_value'*, *reviews_per_month'*.

Three options for removing outliers were available to us, the Interquartile range (IQR) method, Winsorize and Log transform. Log transformation reduces the skewness of data and tries to make it normal. However, we had zero or negative values contained within some of our features. With winsorizing, any value of a variable above or below a percentile k on each side of the variables' distribution is replaced with the value of the k -th percentile itself. For example, 90% winsorization means the replacement of the top 5% and bottom 5% of the data. The top 5% of the data is replaced by the value of the data at the 95th percentile and the value of the bottom 5% of the data is replaced by the value of the data at the 5th percentile. However, winsorization would not be feasible as it could not be applied properly to features in our data and would replace values in the dataset.

For the above reasons, we selected the IQR method. For the IQR method, the third quartile (75th percentile) and first quartile (25th percentile) is subtracted to get the IQR. Any numbers less than the First quartile subtracted from $1.5 \times \text{IQR}$ is considered an outlier and removed; whereas any number greater than the third quartile subtracted from $1.5 \times \text{IQR}$ is considered an outlier and removed. We applied the IQR method using Python, to all the numerical features in our dataset to remove outliers. Figure 5 in the appendix shows outliers in Price and Minimum nights and after removal of outliers, using the describe function. The data frame named 'df4' contained features such as price with outliers, whereas the data frame named 'df' contains features such as price without outliers.

Part II

Feature Engineering

Table 4 in the appendix represents the features in the cleaned dataset that were feature engineered. These features were 'host_since_length' (length of time the host has been a host of an Airbnb), 'labels_host_acceptance_rate' (categorical variable containing levels of host acceptance rate) and 'labels_host_response_rate' (categorical variable containing levels of host response rate). To extract some useful features from our dataset we converted two relevant features (host response rate and host acceptance rate) to categorical features, using binning and labels. We felt the response rate of the host would have a strong influence on both price and review ratings score. Host response rate refers to the host responding to the request of the client or customer before renting Airbnb, during the stay in Airbnb. The requests from clients could be on inquiries on the host listing location and characteristics of the rental space. Since, price fluctuated with neighbourhood or Toronto district, we felt there would be a relationship with the response rate of the host. If there was a higher price for rent, it was likely that the host was responding on time to the inquiries of the customer. A more direct relationship would be between review ratings score and response rate of host. We felt customers would give a higher score on review ratings if the host responded often to their inquiries before the transaction and during the rental period. Host response rate was formatted with % sign removed and was binned into 5 bins and labels were assigned as in very strict, strict, accepting and very accepting. For 'labels_host_response_rate', five bins were also created with labels 'very slow', 'slow', 'fast', 'very fast'.

Another categorical feature that was feature engineered was 'labels_host_acceptance_rate'. The feature 'host acceptance rate' was defined as an 'object' variable in python, with numbers ranging from 100% to 50%, 30% and 0%. We felt that in locations of Toronto and for listings where hosts showed a low acceptance rate there would be a strong relationship with both price of rent and review ratings score. Perhaps with a high acceptance rate, the price of rent was low which may reflect the quality of the rental space. In terms of review ratings score, perhaps higher review ratings score was provided to hosts who had a high acceptance rate.

Finally, we thought that a factor that a customer would consider when choosing to rent an Airbnb in Toronto would be the length of time the host has been a host of an Airbnb. Perhaps the customer considered the experience of the host as measured by the number of years the individual has been an Airbnb host as an important factor in renting the host's Airbnb. We split the date into year under the column "Host_since_year" and month using Python and then subtracted the current year being 2021 from "Host_since_year" to get the new feature "Host_length". *Figure 6* in the Appendix displays a sample of the values for the new features.

Imputation of missing values in Numerical features

There were NaN values present in some of our numerical features. We have highlighted these features in *Table 1*. Using Python code, we imported the library called IterativeImputer to enable an iterative imputer from sklearn to impute NaN values in the numerical features in the Airbnb dataset, with the prediction method. The Iterative Imputer models the missing values based on a prediction compared to the other features in the data set. i.e., known variables are used as a train set and used to predict the test (missing) rows. Essentially missing values are predicted. Moreover, there is a sequential imputation of each feature which allows prior imputed values to be used as part of a model in predicting subsequent features (Reference: Machine learning mastery website). The profile of the cleaned dataset with features is shown in *Table 4* of the Appendix.

Dummy coding categorical variables

In anticipation of the implementation of Machine learning algorithms as in Regression with price being the target variable and a possible implementation of clustering of the data, categorical features were dummy coded using the get_dummies function in Python Pandas. The complete set of features including features which are dummy coded are presented in *Table 4* in the Appendix.

Part III

Exploratory Data Analysis

Using Tableau, we conducted an exploratory data analysis of both the original and cleansed Airbnb datasets.

EDA of original dataset

Before conducting the explanatory analysis of the original dataset, we needed to assess which variables or features were most correlated with our target variable being price and the other target variable being review scores rating. Our correlation plot generated in Python shown in the Figure 2 of the Appendix, showed ‘accommodates’ (i.e., capacity limit for property) ($r = .26$), ‘bed’ (i.e., number of beds in property) ($r = .22$) and ‘bedrooms’ (i.e., number of bedrooms in property) ($r = .24$) to be the three features that were more correlated to ‘price’ than other features. What was also interesting was that ‘review ratings score’ had almost no relationship with ‘price’ ($r = .02$).

Our EDA of the original data also showed there were missing values or NaN values for most of the features and there were outliers for most of the features, even in the target variables as in Price or Review Ratings Score. Our EDA in Tableau showed why Price was most strongly correlated to ‘accommodates’, ‘bed’ and ‘bedrooms’. The three figures shown in Figure 7 of the Appendix show scatter plots of ‘price’ with the features ‘accommodates’, ‘beds’ and ‘bedrooms’.

However, correlation does not measure the relationship between numerical and categorical variables. As a result, several categorical features could relate to average price or review ratings score. One feature that could also relate to price or review ratings score was the location of the Airbnb listing. When we plotted in Tableau the listings by neighbourhood, a bird’s eye view gave us the impression that in the district called “Old City Toronto” there were more listings compared to other districts such as Scarborough and York as can be seen from *Figure B* in the appendix, combined with the fact that price seemed higher for listings in location close to waterfront neighborhood in Old City of Toronto than areas further away from the waterfront neighborhood, such as North York and Etobicoke. As a result, our first business query was: Does

Average price and rating vary by the location of the Airbnb listing? As can be seen from the Figure below, price per day was much higher for listings in the Old City of Toronto (\$158) compared to other districts. Lowest average price per listing was for listings in York (\$92) and Scarborough (\$96) which were located furthest from Old City of Toronto. Interestingly, review scores rating did not vary by location of listing, as shown from the figure below.



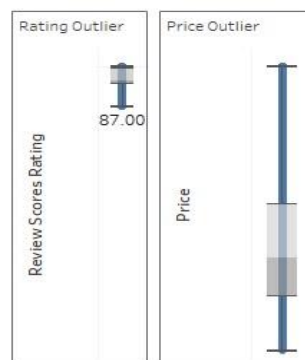
Our second business query was: Does average price and total host listings also vary by the room type of the property and by the district in Toronto? From our EDA of the unclean data, we found that Average price of listing in a district could also depend on the type of room (i.e., Entire home/apartment, Hotel room, Private room and Share room) and in the district in Toronto. As seen from the Figure 8 in the Appendix, there were more listings of Entire home/apartment in the Old City of Toronto, North York and Scarborough compared to other districts. Moreover, the Old City of Toronto seemed to have listings which offered a more diverse distribution of room type. Finally, the average price of Entire home/apartment was the highest in the Old City of Toronto and lower in York, Etobicoke, and East York.

Figure 9 in the Appendix below displays the Tableau dashboard with the uncleaned data being the data source and shows Airbnb house/room ratings between August 2009 to September 2020 in the Greater Toronto Area. Using the cleansed data set, total listings count is at 19,339. Old City of Toronto holds the

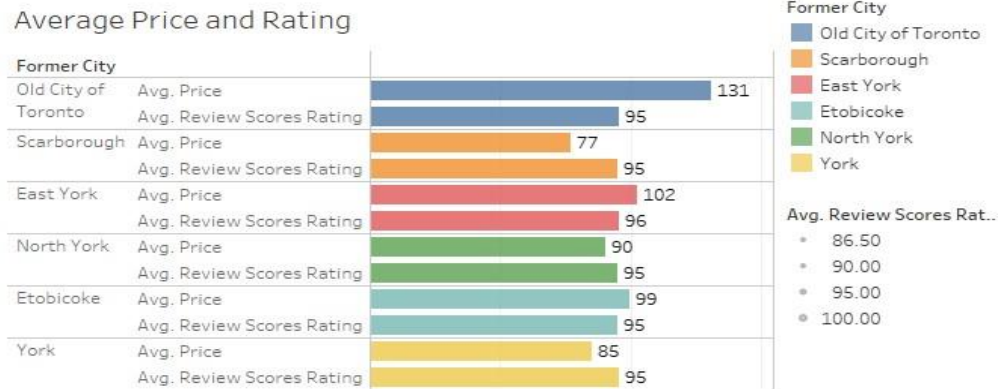
highest average price of \$158 per night and highest average review score at 95. Outliers are shown in the rating and price outlier boxplots.

Visualization from Tableau of Clean data

We also did an EDA of the cleaned Airbnb dataset; however, the details of the EDA will be provided in deliverable 4 for which dashboards of the cleaned dataset will be presented. The main purpose of showing a brief EDA of the clean data in this report was to provide evidence that missing values and outliers in most of the features were removed using data cleaning processes such as imputation in Python. For example, there were no outliers for both review ratings score and price after the data cleaning process as evidenced by the figure below from visualization in Tableau.



Interestingly, the cleaned data did show that the general patterns present in the original dataset remained the same even after data cleaning. For example, the average price for a host Airbnb listing was still highest in Old City of Toronto compared to other districts of Toronto, as can be seen from the figure below.



We also did a visualization of the relationship between our two features: host response rate and host acceptance rate engineered using Python using a binning process. As can be seen from **Figure 10** in the Appendix, host acceptance rate does vary by district in Toronto. In addition, average price does seem to relate to host acceptance rate in certain districts of Toronto, but not other districts such as the Old City of Toronto.

Conclusion

In summary, our correlation plot did not show strong correlations between ‘price’ and other features in our Airbnb data set. However, when exploring the dataset, we noticed that price of rental per day fluctuated with location of the listing in Toronto. Moreover, there were several missing values and outliers in most of the features as noted from our EDA in Python and Tableau. We conducted a detailed data cleaning process using Python to remove irrelevant features, features with missing values more than the 40% threshold, impute missing values with Sklearn Simple and Iterative imputers, create new features and dummy code the cleaned data set. We also performed a brief EDA of the cleaned data exported from Python to confirm that missing values were imputed, outliers were removed and that there were some general patterns in the preserved such as fluctuation of price by location in Toronto.

In conclusion, we believe that we have prepared a clean Toronto Airbnb dataset and our use of EDA in Tableau and Python will enable stakeholders to understand why essential steps characteristic of data cleaning such as the removal of features and imputation took place. The next step as part of Deliverable

3 report will be to apply a suitable machine learning model such as regression to the clean dataset to identify features that predict price per day of a host listing. Importantly, we will determine using an appropriate machine learning model which features predict whether the customer rents an Airbnb listing for a given rental price.

References

<https://www.kaggle.com/robinkongninglo/toronto-airbnb-dataset>

<https://machinelearningmastery.com/iterative-imputation-for-missing-values-in-machine-learning/>

<https://en.wikipedia.org/wiki/Toronto>

Appendix

Figure A. Listings by Neighborhood in Toronto as displayed in Tableau

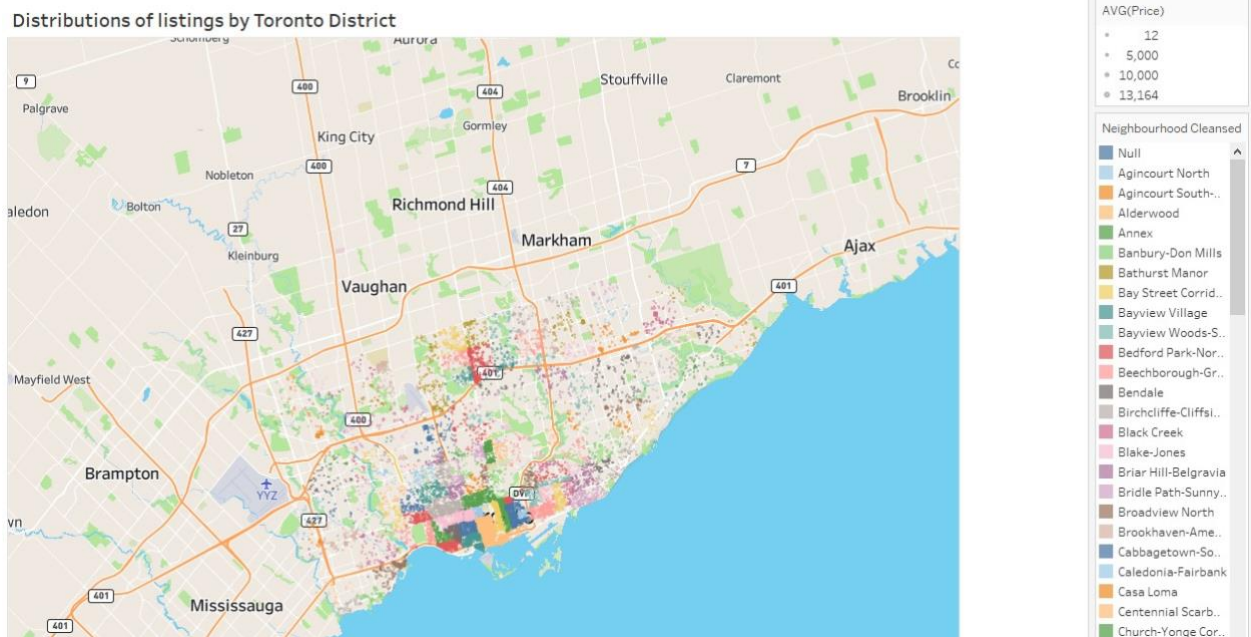


Figure B. Listings by District in Toronto

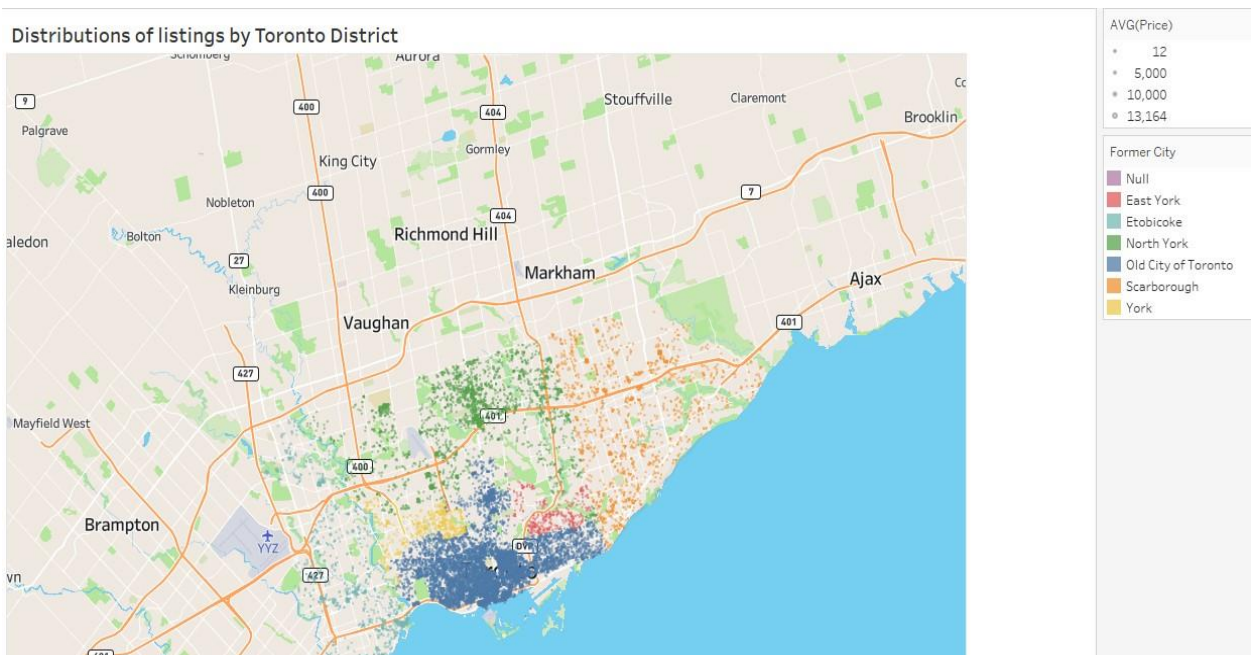


Table 1. Data dictionary showing the fields or features in the AirBnb Data Set with description and the type of feature is stated.

In total there were 73 features in the original dataset that were not cleaned. A colour scheme was implemented to show distribution of features in the unclean or original dataset that were irrelevant, had more than 40% missing values, were target features, contained redundant data or were relevant features.

Legend:

Orange-Irrelevant features

Blue-Features with more than 40% NaN values

Red-Target features

Dark Magenta-Redundant features

Purple-Relevant features

| <u>Number of Variables</u> | <u>Column Name</u> | <u>Description</u> | <u>Feature type(e.g., Numeric, String)</u> |
|----------------------------|-----------------------|--|--|
| | id | Unique listing id (Primary Key) | Numeric |
| 1 | listing_url | Link to the rental property listing on Airbnb | String |
| 2 | scrape_id | Identifier for scraper | Numeric |
| 3 | last_scraped | Last date listing was scraped | Numeric |
| 4 | name | Title of Posting | String |
| 5 | description | Description of Posting | String |
| 6 | neighborhood_overview | Overview of neighborhood | String |
| 7 | picture_url | Link to the main vacation rental listing image on Airbnb | Image |
| 8 | host_id | Unique id for each host | Numeric |
| 9 | host_url | Link to the host profile on Airbnb | String |
| 10 | host_name | Host Name | String |
| 11 | host_since | Date an individual became a host | Numeric |
| 12 | host_location | Location of Listing | String |
| 13 | host_about | Description of host - relationship status, interests and hobbies | String |

| | | | |
|----|-------------------------------------|--|---------|
| 14 | <i>host_response_time</i> | Time to respond to customer booking inquiry; ranges from 1hour to few days | Numeric |
| 15 | <i>host_response_rate</i> | Ranges from 0% to 100% for reply to booking inquiries | Numeric |
| 16 | <i>host_acceptance_rate</i> | Ranges from 0% to 100% response for acceptance of booking | Numeric |
| 17 | <i>host_is_superhost</i> | Is either t(true) or f(false) | Boolean |
| 18 | <i>host_thumbnail_url</i> | Host thumbnail URL | Image |
| 19 | <i>host_picture_url</i> | Host Picture URL | Image |
| 20 | <i>host_neighbourhood</i> | Description of neighborhood listing | String |
| 21 | <i>host_listings_count</i> | Current number of host listings | Numeric |
| 22 | <i>host_total_listings_count</i> | Total number of listings made by host | Numeric |
| 23 | <i>host_verifications</i> | Identifies how the host has completed the identity verification process | String |
| 24 | <i>host_has_profile_pic</i> | Host Profile Pic | Image |
| 25 | <i>host_identity_verified</i> | Identifies if the host has completed the verification process by indicating true or false | Boolean |
| 26 | <i>neighbourhood</i> | Specific location of Toronto area | String |
| 27 | <i>neighbourhood_cleansed</i> | Represents one of boroughs in Toronto in which a listing resides | String |
| 28 | <i>neighbourhood_group_cleansed</i> | No values (N/A) | N/A |
| 29 | <i>latitude</i> | The angular distance of a location or object north or south of the Earth's celestial equator | String |
| 30 | <i>longitude</i> | The angular distance of a location or object east or west of the meridian | String |
| 31 | <i>property_type</i> | Type of property (e.g., Entire house) | String |
| 32 | <i>room_type</i> | Specific type of room (e.g., Entire home/apt) | String |
| 33 | <i>accommodates</i> | How many people can stay (e.g., 6) | Numeric |
| 34 | <i>bathrooms</i> | No values, NA | N/A |
| 35 | <i>bathrooms_text</i> | Number of bathrooms in property | Numeric |
| 36 | <i>bedrooms</i> | Number of bedrooms in property | Numeric |

| | | | |
|----|------------------------|---|---------|
| 37 | beds | Number of beds in property | Numeric |
| 38 | amenities | List of amenities such as shampoo available in property | String |
| 39 | price | Price of stay per night | Numeric |
| 40 | minimum_nights | Minimum nights can be booked by same individual | Numeric |
| 41 | maximum_nights | Maximum nights can be booked by same individual | Numeric |
| 42 | minimum_minimum_nights | Same values as Minimum nights | Numeric |
| 43 | maximum_minimum_nights | Same values as Maximum_nights | Numeric |
| 44 | minimum_maximum_nights | Same values as Maximum_nights | Numeric |
| 45 | maximum_maximum_nights | Same values as Maximum_nights | Numeric |
| 46 | minimum_nights_avg_ntm | Average minimum nights can be booked by same individual | Numeric |
| 47 | maximum_nights_avg_ntm | Average maximum nights can be booked by same individual | Numeric |
| 48 | calendar_updated | No values (N/A) | N/A |
| 49 | has_availability | True or False if listing is available | Boolean |
| 50 | availability_30 | Availability of Property in 30 days | Numeric |
| 51 | availability_60 | Availability of Property in 60 days | Numeric |
| 52 | availability_90 | Availability of Property in 90 days | Numeric |
| 53 | availability_365 | Availability of Property in 365 days | Numeric |
| 54 | calendar_last_scraped | Date last scraped | Numeric |
| 55 | number_of_reviews | Total number of reviews that a listing has received from customers | Numeric |
| 56 | number_of_reviews_ltm | The number of reviews that a listing has received last twelve month | Numeric |
| 57 | number_of_reviews_l30d | The number of reviews that a listing has received per 130 days | Numeric |
| 58 | first_review | Date of first review by customer | Numeric |
| 59 | last_review | Date of last review by customer | Numeric |

| | | | |
|----|--|---|---------|
| 60 | review_scores_rating | Customer-provided score rating (0% to 100%); A customer-provided review score attributed to a listing based on overall experience and satisfaction | Numeric |
| 61 | review_scores_accuracy | Accuracy of review scores (0 to 10) | Numeric |
| 62 | review_scores_cleanliness | Cleanliness score (0 to 10) | Numeric |
| 63 | review_scores_checkin | Over-all check in score (0 to 10) | Numeric |
| 64 | review_scores_communication | Score on communication with host (0 to 10) | Numeric |
| 65 | review_scores_location | Score on location based on factors such as nearby transportation, noise level (0 to 10) | Numeric |
| 66 | review_scores_value | Over- all value/quality/experience (0 to 10) | Numeric |
| 67 | license | No values (N/A) | N/A |
| 68 | instant_bookable | True or False if customer can instantly book | Boolean |
| 69 | calculated_host_listings_count | Calculated number of listings by host | Numeric |
| 70 | calculated_host_listings_count_entire_homes | Number of host listings which are entire homes | Numeric |
| 71 | calculated_host_listings_count_private_rooms | Number of host listings which are private rooms | Numeric |
| 72 | calculated_host_listings_count_shared_rooms | Number of host listings which are shared homes | Numeric |
| 73 | reviews_per_month | Number of Customer reviews of the host accommodation or accommodations per month | Numeric |

| | price | review_scores_rating | number_of_reviews | accommodates | bedrooms | beds | availability_30 | number_of_reviews | reviews_per_month | calculated_host_listings_count | calculated_host_listings_count_private_rooms | calculated_host_listings_count_shared_rooms | host_total_listings_count | host_listings_count | minimum_nights | maximum_nights |
|--|---------|----------------------|-------------------|--------------|----------|--------|-----------------|-------------------|-------------------|--------------------------------|--|---|---------------------------|---------------------|----------------|----------------|
| price | 1 | 0.02 | -0.016 | 0.26 | 0.24 | 0.22 | 0.053 | -0.016 | 0.0056 | 0.00032 | 0.079 | -0.033 | 0.049 | 0.0049 | 0.031 | 0.00055 |
| review_scores_rating | 0.02 | 1 | 0.081 | 0.0062 | 0.0057 | 0.0057 | 0.071 | 0.081 | 0.082 | -0.099 | -0.12 | -0.051 | -0.083 | -0.083 | 0.017 | 0.0058 |
| number_of_reviews | -0.016 | 0.081 | 1 | 0.078 | -0.019 | 0.032 | 0.039 | 1 | 0.76 | -0.043 | -0.063 | -0.033 | -0.031 | -0.031 | -0.031 | 0.0094 |
| accommodates | 0.26 | 0.0062 | 0.078 | 1 | 0.77 | 0.8 | 0.0064 | 0.078 | 0.11 | 0.041 | -0.19 | -0.074 | 0.03 | 0.03 | -0.003 | 0.0033 |
| bedrooms | 0.24 | 0.005 | -0.019 | 0.77 | 1 | 0.8 | -0.015 | -0.019 | 0.008 | 0.037 | -0.14 | -0.043 | 0.0007 | 0.0074 | 0.014 | 0.0038 |
| beds | 0.22 | 0.005 | 0.032 | 0.8 | 0.8 | 1 | 0.0063 | 0.032 | 0.059 | 0.027 | -0.14 | -0.032 | 0.018 | 0.018 | -0.0012 | 0.0024 |
| availability_30 | -0.053 | -0.071 | 0.039 | 0.0064 | -0.015 | 0.0063 | 1 | 0.039 | 0.14 | 0.14 | 0.098 | 0.083 | 0.12 | 0.12 | 0.095 | 0.01 |
| number_of_reviews | -0.016 | 0.081 | 1 | 0.078 | -0.019 | 0.032 | 0.039 | 1 | 0.76 | -0.043 | -0.063 | -0.033 | -0.031 | -0.031 | -0.031 | 0.0094 |
| reviews_per_month | -0.0056 | 0.082 | 0.76 | 0.11 | -0.0086 | 0.059 | 0.14 | 0.76 | 1 | -0.023 | -0.06 | -0.027 | -0.022 | -0.022 | -0.053 | 0.0026 |
| calculated_host_listings_count | 0.00032 | 0.099 | -0.043 | 0.041 | 0.0037 | 0.027 | 0.14 | -0.043 | -0.023 | 1 | 0.34 | 0.32 | 0.76 | 0.76 | -0.011 | 0.0021 |
| calculated_host_listings_count_private_rooms | -0.079 | -0.12 | -0.063 | -0.19 | -0.14 | -0.14 | 0.098 | -0.063 | -0.06 | 0.34 | 1 | 0.084 | 0.24 | 0.24 | -0.035 | 0.0032 |
| calculated_host_listings_count_shared_rooms | -0.033 | -0.051 | -0.033 | -0.074 | -0.043 | -0.032 | 0.083 | -0.033 | -0.027 | 0.32 | 0.084 | 1 | 0.17 | 0.17 | -0.013 | 0.0059 |
| host_total_listings_count | -0.0049 | -0.083 | -0.031 | 0.03 | 0.00074 | 0.018 | 0.12 | -0.031 | -0.022 | 0.76 | 0.24 | 0.17 | 1 | 1 | 0.2e-05 | 0.00069 |
| host_listings_count | -0.0049 | -0.083 | -0.031 | 0.03 | 0.00074 | 0.018 | 0.12 | -0.031 | -0.022 | 0.76 | 0.24 | 0.17 | 1 | 1 | 0.2e-05 | 0.00069 |
| minimum_nights | -0.031 | 0.017 | -0.031 | 0.0038 | 0.014 | 0.0012 | 0.095 | -0.031 | -0.053 | -0.011 | -0.035 | -0.0139 | 2e-05 | 2e-05 | 1 | 0.0015 |
| maximum_nights | 0.00055 | 0.0058 | 0.0094 | 0.0033 | 0.0038 | 0.0024 | 0.01 | 0.0094 | 0.0026 | 0.0021 | 0.0032 | 0.00059 | 0.00069 | 0.00069 | 0.0015 | 1 |

Table 2. Number of non-null values by each feature displayed in dataset feature profile displayed in output from python's dataframe.info() function. There were 74 features, 41 numerical(21 integer, 20 float) and 34 categorical (object) features. As can be seen missing values were present in most features. In addition, there were diversity of data types characterizing the features. For example, 'neighbourhood' is a categorical feature (i.e., Object), 'price' is listed as object due to being in currency format, review ratings score which is a numerical feature is of float data type and 'accomodates' which is a numerical feature is of integer data type.

<class 'pandas.core.frame.DataFrame'>

Int64Index: 19343 entries, 0 to 19342

Data columns (total 75 columns):

| # | Column | Non-Null Count | Dtype |
|----|-----------------------|----------------|--------|
| 0 | id | 19343 non-null | int64 |
| 1 | listing_url | 19343 non-null | object |
| 2 | scrape_id | 19343 non-null | int64 |
| 3 | last_scraped | 19343 non-null | object |
| 4 | name | 19342 non-null | object |
| 5 | description | 18623 non-null | object |
| 6 | neighborhood_overview | 12364 non-null | object |
| 7 | picture_url | 19343 non-null | object |
| 8 | host_id | 19343 non-null | int64 |
| 9 | host_url | 19343 non-null | object |
| 10 | host_name | 19339 non-null | object |
| 11 | host_since | 19339 non-null | object |
| 12 | host_location | 19329 non-null | object |

| | |
|--|-------------------------------|
| 13 host_about | 10958 non-null object |
| 14 host_response_time | 11814 non-null object |
| 15 host_response_rate | 11814 non-null object |
| 16 host_acceptance_rate | 13672 non-null object |
| 17 host_is_superhost | 19339 non-null object |
| 18 host_thumbnail_url | 19339 non-null object |
| 19 host_picture_url | 19339 non-null object |
| 20 host_neighbourhood | 15552 non-null object |
| 21 host_listings_count | 19339 non-null float64 |
| 22 host_total_listings_count | 19339 non-null float64 |
| 23 host_verifications | 19343 non-null object |
| 24 host_has_profile_pic | 19339 non-null object |
| 25 host_identity_verified | 19339 non-null object |
| 26 neighbourhood | 12364 non-null object |
| 27 neighbourhood_cleansed | 19343 non-null object |
| 28 neighbourhood_group_cleansed | 0 non-null float64 |
| 29 latitude | 19343 non-null float64 |
| 30 longitude | 19343 non-null float64 |
| 31 property_type | 19343 non-null object |
| 32 room_type | 19343 non-null object |
| 33 accommodates | 19343 non-null int64 |

| | | |
|-----------|-------------------------------|-------------------------------|
| 34 | bathrooms | 0 non-null float64 |
| 35 | bathrooms_text | 19330 non-null object |
| 36 | bedrooms | 17914 non-null float64 |
| 37 | beds | 19147 non-null float64 |
| 38 | amenities | 19343 non-null object |
| 39 | price | 19343 non-null object |
| 40 | minimum_nights | 19343 non-null int64 |
| 41 | maximum_nights | 19343 non-null int64 |
| 42 | minimum_minimum_nights | 19343 non-null int64 |
| 43 | maximum_minimum_nights | 19343 non-null int64 |
| 44 | minimum_maximum_nights | 19343 non-null int64 |
| 45 | maximum_maximum_nights | 19343 non-null int64 |
| 46 | minimum_nights_avg_ntm | 19343 non-null float64 |
| 47 | maximum_nights_avg_ntm | 19343 non-null float64 |
| 48 | calendar_updated | 0 non-null float64 |
| 49 | has_availability | 19343 non-null object |
| 50 | availability_30 | 19343 non-null int64 |
| 51 | availability_60 | 19343 non-null int64 |
| 52 | availability_90 | 19343 non-null int64 |
| 53 | availability_365 | 19343 non-null int64 |
| 54 | calendar_last_scraped | 19343 non-null object |

| | | |
|----|--|------------------------|
| 55 | number_of_reviews | 19343 non-null int64 |
| 56 | number_of_reviews_ltm | 19343 non-null int64 |
| 57 | number_of_reviews_l30d | 19343 non-null int64 |
| 58 | first_review | 15278 non-null object |
| 59 | last_review | 15278 non-null object |
| 60 | review_scores_rating | 15010 non-null float64 |
| 61 | review_scores_accuracy | 14976 non-null float64 |
| 62 | review_scores_cleanliness | 14976 non-null float64 |
| 63 | review_scores_checkin | 14974 non-null float64 |
| 64 | review_scores_communication | 14978 non-null float64 |
| 65 | review_scores_location | 14971 non-null float64 |
| 66 | review_scores_value | 14972 non-null float64 |
| 67 | license | 0 non-null float64 |
| 68 | instant_bookable | 19343 non-null object |
| 69 | calculated_host_listings_count | 19343 non-null int64 |
| 70 | calculated_host_listings_count_entire_homes | 19343 non-null int64 |
| 71 | calculated_host_listings_count_private_rooms | 19343 non-null int64 |
| 72 | calculated_host_listings_count_shared_rooms | 19343 non-null int64 |
| 73 | reviews_per_month | 15278 non-null float64 |
| 74 | former_city | 19343 non-null object |

dtypes: float64(20), int64(21), object(34)

Table 3. Feature dataset profile after imputation of missing values for numerical features

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 19343 entries, 0 to 19342
Data columns (total 37 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   host_total_listings_count                 19339 non-null   float64
1   neighbourhood_cleansed                   19343 non-null   object
2   latitude                                 19343 non-null   float64
3   longitude                                 19343 non-null   float64
4   accommodates                             19343 non-null   int64
5   bedrooms                                 17914 non-null   float64
6   beds                                     19147 non-null   float64
7   price                                    19343 non-null   float64
8   minimum_nights                          19343 non-null   int64
9   maximum_nights                          19343 non-null   int64
10  availability_30                          19343 non-null   int64
11  availability_60                          19343 non-null   int64
12  availability_90                          19343 non-null   int64
13  availability_365                         19343 non-null   int64
14  number_of_reviews                       19343 non-null   int64
15  number_of_reviews_ltm                   19343 non-null   int64
16  number_of_reviews_l30d                  19343 non-null   int64
17  review_scores_rating                    15010 non-null   float64
18  review_scores_accuracy                  14976 non-null   float64
19  review_scores_cleanliness               14976 non-null   float64
20  review_scores_checkin                   14974 non-null   float64
21  review_scores_communication             14978 non-null   float64
22  review_scores_location                  14971 non-null   float64
23  review_scores_value                     14972 non-null   float64
24  instant_bookable                       19343 non-null   object
25  calculated_host_listings_count          19343 non-null   int64
26  calculated_host_listings_count_entire_homes 19343 non-null   int64
27  calculated_host_listings_count_private_rooms 19343 non-null   int64
28  calculated_host_listings_count_shared_rooms 19343 non-null   int64
29  reviews_per_month                      15278 non-null   float64
30  former_city                             19343 non-null   object
31  room_type                               19343 non-null   object
32  host_since                              19343 non-null   object
33  host_response_time                      19343 non-null   object
34  host_is_superhost                       19343 non-null   object
35  host_acceptance_rate                    19343 non-null   object
36  host_response_rate                      19343 non-null   object
dtypes: float64(14), int64(14), object(9)
memory usage: 5.6+ MB
```

Figure 4. Outliers in numerical features as shown from Python and in Tableau

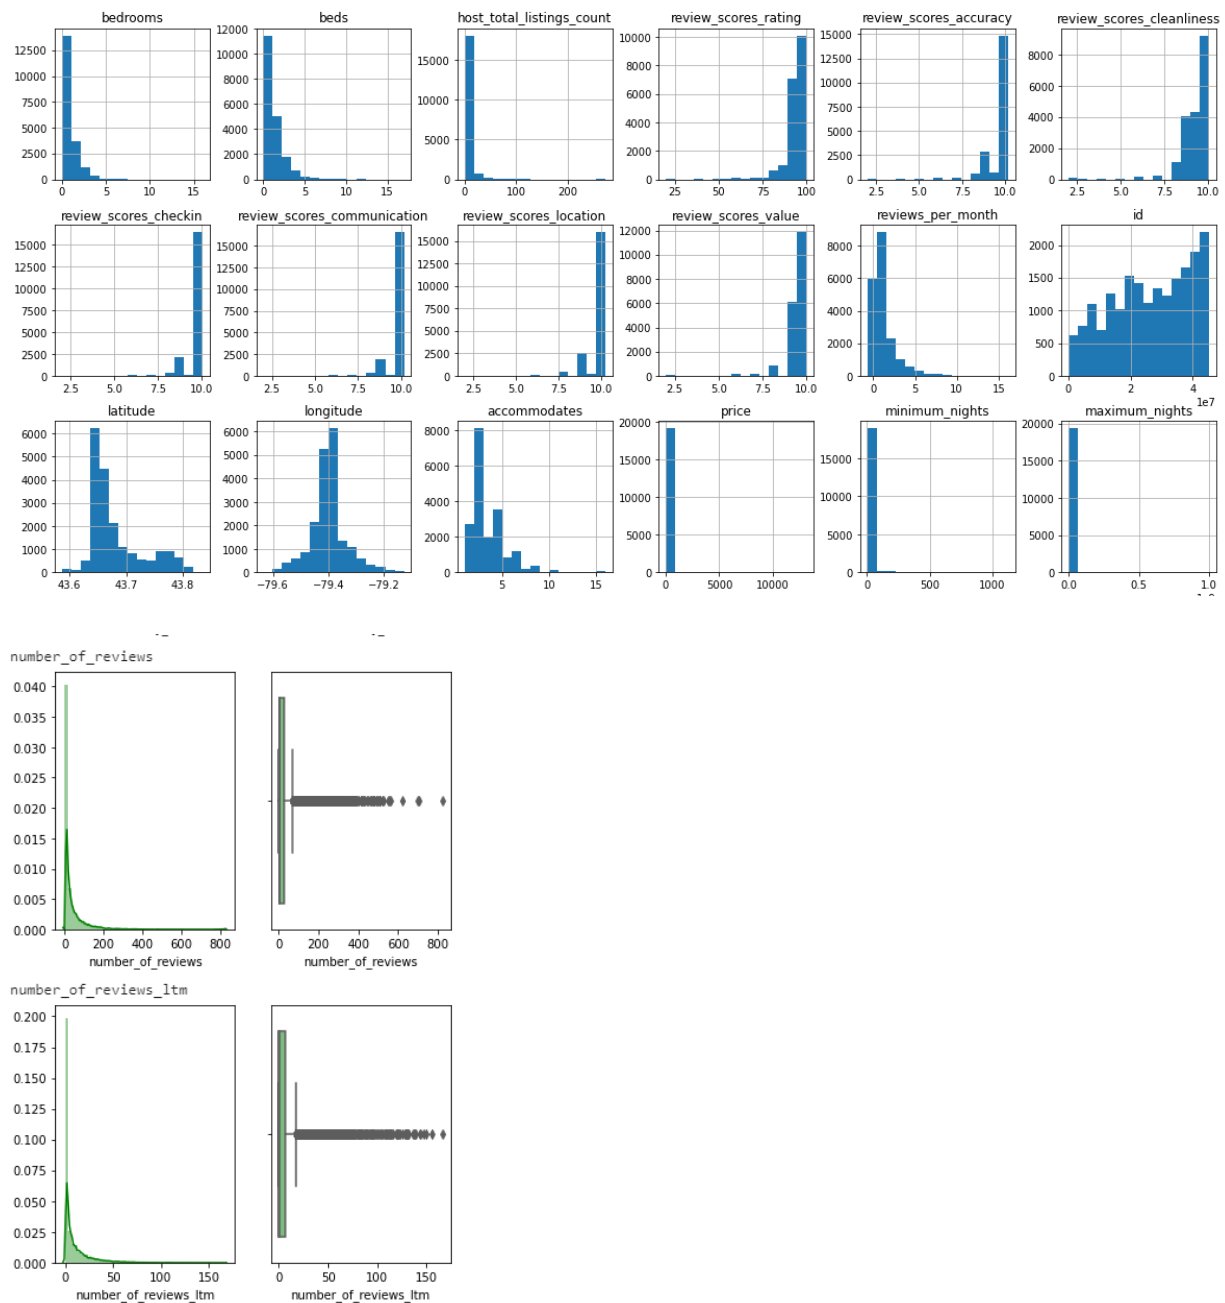


Figure 5. The dataframe named 'df4' contained features such as price with outliers, whereas the data frame named 'df' contains features such as price without outliers.

```
df4[['price', 'review_scores_rating', 'minimum_nights', 'host_total_listings_count']].describe()
```

| | price | review_scores_rating | minimum_nights | host_total_listings_count |
|-------|--------------|----------------------|----------------|---------------------------|
| count | 19343.000000 | 15010.000000 | 19343.000000 | 19339.000000 |
| mean | 141.278116 | 94.304930 | 9.968730 | 5.52283 |
| std | 290.664182 | 8.899373 | 37.022953 | 15.39491 |
| min | 12.000000 | 20.000000 | 1.000000 | 0.000000 |
| 25% | 63.000000 | 93.000000 | 1.000000 | 1.000000 |
| 50% | 100.000000 | 97.000000 | 2.000000 | 1.000000 |
| 75% | 150.000000 | 100.000000 | 5.000000 | 4.000000 |
| max | 13164.000000 | 100.000000 | 1125.000000 | 272.000000 |

```
df[['price', 'review_scores_rating', 'minimum_nights', 'host_total_listings_count']].describe()
```

| | price | review_scores_rating | minimum_nights | host_total_listings_count |
|-------|--------------|----------------------|----------------|---------------------------|
| count | 19343.000000 | 19343.000000 | 19343.000000 | 19343.000000 |
| mean | 117.516891 | 95.257847 | 3.834876 | 2.742685 |
| std | 70.959498 | 4.043789 | 3.598508 | 2.770920 |
| min | 12.000000 | 86.500000 | 1.000000 | 0.000000 |
| 25% | 63.000000 | 94.000000 | 1.000000 | 1.000000 |
| 50% | 100.000000 | 95.000000 | 2.000000 | 1.000000 |
| 75% | 150.000000 | 99.000000 | 5.000000 | 4.000000 |
| max | 280.500000 | 100.000000 | 11.000000 | 8.500000 |

Figure 6. Display of sample values in new features

```
df[['host_acceptance_rate', 'labels_host_acceptance_rate', 'host_response_rate', 'labels_h
```

| | host_acceptance_rate | labels_host_acceptance_rate | host_response_rate | labels_host_response_rate | host_since_year |
|-------|----------------------|-----------------------------|--------------------|---------------------------|-----------------|
| 0 | 100 | Very Accepting | 100 | Very high | 2008 |
| 1 | 100 | Very Accepting | 100 | Very high | 2011 |
| 2 | 100 | Very Accepting | 100 | Very high | 2012 |
| 3 | 100 | Very Accepting | 100 | Very high | 2012 |
| 4 | 100 | Very Accepting | 100 | Very high | 2012 |
| ... | ... | ... | ... | ... | ... |
| 19338 | 100 | Very Accepting | 100 | Very high | 2019 |
| 19339 | 100 | Very Accepting | 100 | Very high | 2020 |
| 19340 | 94 | Very Accepting | 100 | Very high | 2017 |
| 19341 | 75 | Accepting | 100 | Very high | 2019 |
| 19342 | 100 | Very Accepting | 100 | Very high | 2020 |

19343 rows × 5 columns

Table 4. Feature dataset profile of the Cleaned dataset that will be used for deliverables; Part 3 and Part 4

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 19343 entries, 0 to 19342
Data columns (total 41 columns):
#   Column                                     Non-Null Count  Dtype
---  -
0   bedrooms                                 19343 non-null  float64
1   beds                                    19343 non-null  float64
2   host_total_listings_count               19343 non-null  float64
3   review_scores_rating                    19343 non-null  float64
4   review_scores_accuracy                  19343 non-null  float64
5   review_scores_cleanliness               19343 non-null  float64
6   review_scores_checkin                   19343 non-null  float64
7   review_scores_communication             19343 non-null  float64
8   review_scores_location                  19343 non-null  float64
9   review_scores_value                     19343 non-null  float64
10  reviews_per_month                       19343 non-null  float64
11  neighbourhood_cleansed                   19343 non-null  object
12  latitude                                 19343 non-null  float64
13  longitude                                19343 non-null  float64
14  accommodates                             19343 non-null  int64
15  price                                    19343 non-null  float64
16  minimum_nights                           19343 non-null  int64
17  maximum_nights                           19343 non-null  int64
18  availability_30                           19343 non-null  int64
19  availability_60                           19343 non-null  int64
20  availability_90                           19343 non-null  int64
21  availability_365                           19343 non-null  int64
22  number_of_reviews                         19343 non-null  int64
23  number_of_reviews_ltm                     19343 non-null  int64
24  number_of_reviews_l30d                     19343 non-null  int64
25  instant_bookable                         19343 non-null  object
26  calculated_host_listings_count            19343 non-null  int64
27  calculated_host_listings_count_entire_homes  19343 non-null  int64
28  calculated_host_listings_count_private_rooms  19343 non-null  int64
29  calculated_host_listings_count_shared_rooms  19343 non-null  int64
30  former_city                             19343 non-null  object
31  room_type                               19343 non-null  object
32  host_since                             19343 non-null  object
33  host_response_time                       19343 non-null  object
34  host_is_superhost                       19343 non-null  object
35  host_acceptance_rate                     19343 non-null  int64
36  host_response_rate                       19343 non-null  int64
37  labels_host_acceptance_rate               19343 non-null  category
38  labels_host_response_rate                 19343 non-null  category
39  host_since_year                           19343 non-null  int64
40  host_length                             19343 non-null  int64
dtypes: category(2), float64(14), int64(18), object(7)
memory usage: 5.8+ MB
```

Figure 7.

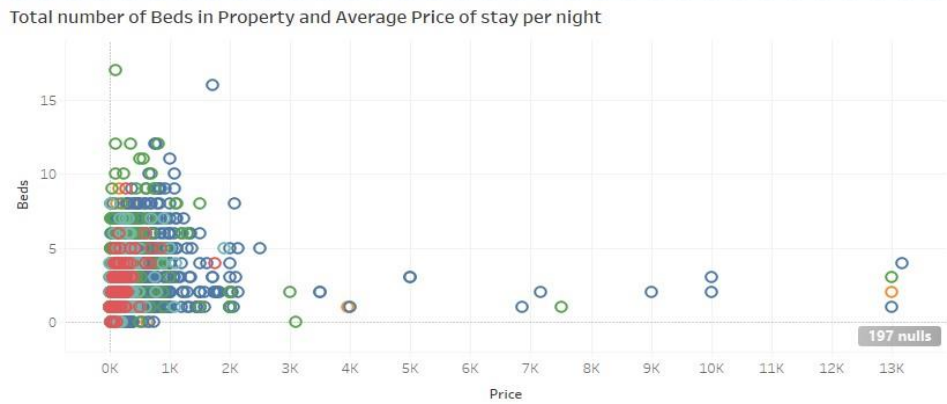


Figure 8. Room type and Average price per day of listing by District in Toronto

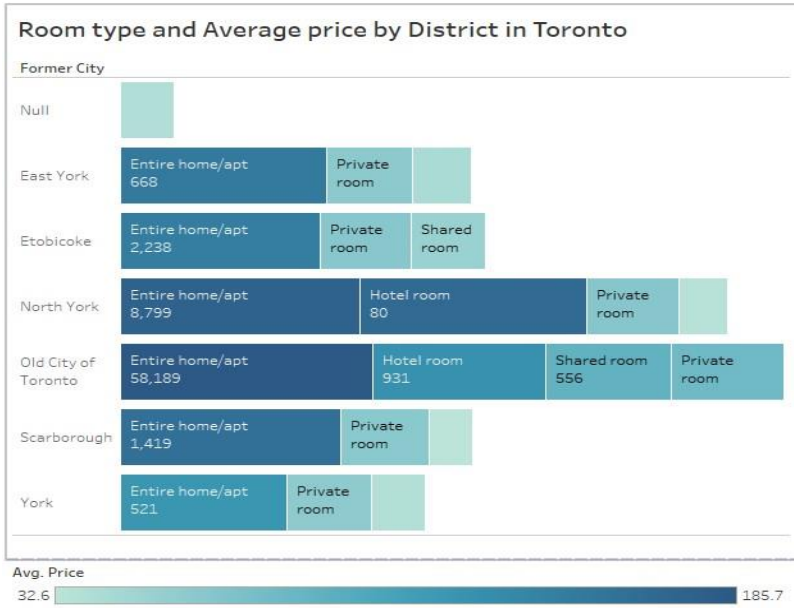


Figure 9. Tableau Dashboard with the uncleaned data being the data source

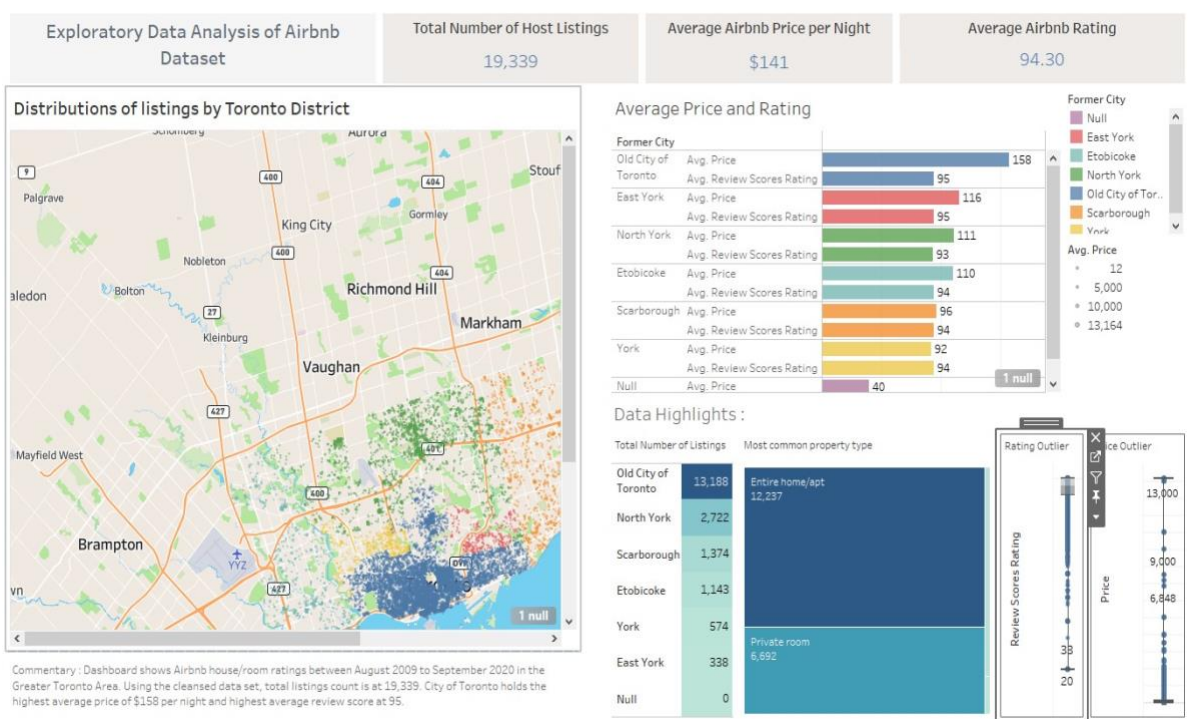


Figure 10. Average price and Host Acceptance rate by District in Toronto

