


[Open in app](#)[Get started](#)

Published in Towards Data Science

You have **1** free member-only story left this month. [Sign up for Medium and get an extra one](#)



Thu Vu

[Follow](#)

Jan 16, 2020 · 5 min read · · Listen




Save



Blazingly Fast Data Wrangling With R data.table

Who has time to do data science with slow code?



[Open in app](#)[Get started](#)

Update March 2020: You can find a very interesting comparison between `data.table` and `dplyr` [here](#).

Introduction

I have recently noticed that every R script I wrote starts with `library(data.table)` . And that seems a very compelling reason for me to write a post about it.

You may have seen that as the machine learning community is developing, Python has gained enormous popularity and as a result, the Pandas library has automatically become a staple for a lot of people.

However, if I had a choice, I would definitely prefer using R [data.table](#) for relatively large dataset data manipulation, in-depth exploration and ad-hoc analyses. Why? Because it's so fast, elegant and beautiful (sorry if I got too enthusiastic!). But hey, working with data means you'll have to turn the code upside down, mould it, clean it for maybe... uhm... a thousand times before you can actually build a machine learning model. In fact, data pre-processing usually takes 80–90% the time of a data science project. That's why I can't afford slow and inefficient code (and I don't think anyone should, as a data professional).

So let's dive into what `data.table` is and why [many people](#) have become big fans of it.

1. So what the heck is `data.table`?

`data.table` [package](#) is an extension of `data.frame` package in R. It is widely used for fast aggregation of large datasets, low latency add/update/remove of columns, quicker ordered joins, and a fast file reader.

That sounds good, right? You may think it's difficult to pick up, but actually a



[Open in app](#)[Get started](#)

2. Data.table is extremely fast

From my own experience, working with fast code can really improve the thinking flow in the data analysis process. Speed also very important in a data science project, in which you usually have to quickly prototype an idea.

When it comes to speed, `data.table` puts all other packages in Python and many other languages to shame. This is shown in this [benchmark](#), which compares tools from R, Python and Julia. To do five data manipulations on a 50GB dataset, `data.table` only took on average 123s, while Spark took 381s, (py)datatable took 712s, and `pandas` could not do the task due to out of memory.

One of the most powerful functions in `data.table` package is `fread()`, which imports data similarly to what `read.csv()` does. But it's optimized and much much faster. Let's look at this example:

```
require("microbenchmark")
res <- microbenchmark(
  read.csv = read.csv(url("https://archive.ics.uci.edu/ml/machine-
learning-databases/adult/adult.data"), header=FALSE),
  fread = data.table::fread("https://archive.ics.uci.edu/ml/machine-
learning-databases/adult/adult.data", header=FALSE),
  times = 10)
res
```

The results will show that on average, `fread()` will be 3-4 times faster than `read.csv()` function.

3. It is intuitive and elegant

Almost every data manipulation with `data.table` will look like this:



[Open in app](#)[Get started](#)

Source: <https://github.com/Rdatatable/data.table/wiki>

As a result, the code you write will be very consistent and easy to read. Let's take the US Census Income dataset for illustration purposes:

```
dt <- fread("https://archive.ics.uci.edu/ml/machine-learning-  
databases/adult/adult.data", header=FALSE)  
names(dt) <- c("age", "workclass", "fnlwgt", "education",  
"education_num", "marital_status", "occupation", "relationship",  
"race", "sex", "capital_gain", "capital_loss", "hours_per_week",  
"native_country", "label")
```

In the below examples, I'll compare the code in base R, Python and data.table, so that you can easily compare them:

1. To compute the *average age of all "Tech-support" workers*:

> in base R you'd probably write something like this:

```
mean(dt[dt$occupation == 'Tech-support', 'age'])
```

> in Python:



[Open in app](#)[Get started](#)

```
dt[occupation == 'Tech-support', mean(age)]
```

As you can see in this simple example, `data.table` removes all redundancy of repeating `dt` all the times, compared to Python and base R. This in turn reduces the chance of making typo mistakes (remember the coding principle DRY — Don't repeat yourself!)

2. To *aggregate age by occupation* for all male workers:

> in **base R** you'd probably write:

```
aggregate(age ~ occupation, dt[dt$sex == 'Male', ], sum)
```

> in **Python**:

```
dt[dt.sex == 'Male'].groupby('occupation')['age'].sum()
```

> **vs. in data.table**:

```
dt[sex == 'Male', .(sum(age)), by = occupation]
```

The `by =` term defines which column(s) you want to aggregate your data on. This `data.table` syntax may seem a little intimidating at first, but once you get used to it you'll never bother typing “`groupby(...)`” again.

3. To *conditionally modify* values in a column, for example to increase the `age` of all



[Open in app](#)[Get started](#)

```
dt$age[dt$occupation == 'Tech-support'] <- dt$age[dt$occupation ==  
'Tech-support'] + 5
```

> in **Python** (there are several alternatives that are equally long):

```
mask = dt['occupation'] == 'Tech-support'  
dt.loc[mask, 'age'] = dt.loc[mask, 'age'] + 5
```

or using `np.where`:

```
dt['age'] = np.where(dt['occupation'] == 'Tech-support', dt.age + 5,  
dt.age)
```

> vs. in **data.table**:

```
dt[occupation == 'Tech-support', age := age + 5]  
  
# and the ifelse function does just as nicely:  
dt[, age := ifelse(occupation == 'Tech-support', age + 5, age)]
```

It's almost like a magic, isn't it? No more cumbersome repetitive code, and it keeps yourself DRY. You may have noticed the strange operator `:=` in the `data.table` syntax. This operator is used to assign new values to an existing column, just like using the argument `inplace=True` in Python.

4. *Renaming columns* is a breeze in `data.table`. If I want to rename `occupation` column as `job`:



[Open in app](#)[Get started](#)

```
colnames(dt)[which(names(dt) == "occupation")] <- "job"
```

> and in **Python**:

```
dt = dt.rename(columns={'occupation': 'job'})
```

> vs. in **data.table**:

```
setnames(dt, 'occupation', 'job')
```

5. What about *applying a function* to several columns? Suppose you want to multiply `capital_gain` and `capital_loss` by 1000:

> in **base R**:

```
apply(dt[,c('capital_gain', 'capital_loss')], 2, function(x) x*1000)
```

> in **Python**:

```
dt[['capital_gain', 'capital_loss']].apply(lambda x : x*1000)
```

> vs. in **data.table**:

```
dt[, lapply(.SD, function(x) x*1000), .SDcols = c("capital_gain",  
"capital_loss")]
```



[Open in app](#)[Get started](#)

preserved this way.

Conclusion:

From a few simple illustrations above, one can see that the code in R `data.table` is in many cases faster, cleaner and more efficient than in base R and Python. The form of `data.table` code is very consistent. You only need to remember:

```
DT[i, j, by]
```

As an additional note, the data subsetting with `i` by *keying* a `data.table` even allows faster subsets, joins and sorts, which you can read more about in this [documentation](#), or this very useful



164



4

Thank you for reading. If you like this post, I would write more about how to do advanced data wrangling with R and Python in the future posts.

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

By signing up, you will create a Medium account if you don't already have one. Review our [Privacy Policy](#) for more information about our privacy practices.

[Get this newsletter](#)



Open in app

Get started

Get the Medium app

